

Lecture 12

Seriation/Clustering

Breitzman 4/24/2017

Introduction

- We haven't done much with clustering in Data Mining 2.
- In Data Mining 1 we talked about k-means clustering and hierarchical clustering
- Tonight we'll talk about seriation.
- Seriation can be used with clustering to clean up the results or if we're lucky...
 - It can be used by itself for clustering
- Go to <http://www.atgc-montpellier.fr/permutmatrix/manual/SeriationCorps.htm>


Simple Example Stolen from Previous Website

Simple example n°2

Step 1 The smallest dissimilarity in D is 2,58. This value is already near the diagonal. No movement is necessary at this step.

| D | 3 | 4 | 1 | 5 | 2 |
|---|------|------|-------|------|------|
| 3 | 0,00 | | | | |
| 4 | 2,58 | 0,00 | | | |
| 1 | 5,33 | 7,90 | 0,00 | | |
| 5 | 5,22 | 2,65 | 10,55 | 0,00 | |
| 2 | 2,73 | 5,29 | 2,61 | 7,94 | 0,00 |

Graphical representation of D




Graphical data matrix




Step 2 The dissimilarity between items 1 and 2 is now the smallest dissimilarity available.
 $D_{min} = 2,61$.
To put this value against the diagonal, we swap items 2 and 5.

| D | 3 | 4 | 1 | 5 | 2 |
|---|------|------|-------|------|------|
| 3 | 0,00 | | | | |
| 4 | 2,58 | 0,00 | | | |
| 1 | 5,33 | 7,90 | 0,00 | | |
| 5 | 5,22 | 2,65 | 10,55 | 0,00 | |
| 2 | 2,73 | 5,29 | 2,61 | 7,94 | 0,00 |



Step 3 At this step, the dissimilarity between items 4 and 5 is the currently smallest dissimilarity. Items 4 and 3 were pooled at the first step and we are not allowed to modify this pair. We must put item 5 between items 4 and 1. At the end of this step, items 3, 4 and 5 form a "fragment" of three elements.

| D | 3 | 4 | 1 | 2 | 5 |
|---|------|------|-------|------|------|
| 3 | 0,00 | | | | |
| 4 | 2,58 | 0,00 | | | |
| 1 | 5,33 | 7,90 | 0,00 | | |
| 2 | 2,73 | 5,29 | 2,61 | 0,00 | |
| 5 | 5,22 | 2,65 | 10,55 | 7,94 | 0,00 |




Step 4 For the final step, the dissimilarity between items 2 and 3 is the smallest dissimilarity available. To put this value against the diagonal, we swap fragments {1,2} and {3,4,5}.

| D | 3 | 4 | 5 | 1 | 2 |
|---|------|------|-------|------|------|
| 3 | 0,00 | | | | |
| 4 | 2,58 | 0,00 | | | |
| 5 | 5,22 | 2,65 | 0,00 | | |
| 1 | 5,33 | 7,90 | 10,55 | 0,00 | |
| 2 | 2,73 | 5,29 | 7,94 | 2,61 | 0,00 |





End At the end of this algorithm, small dissimilarities are spread along the main diagonal of D .
Close elements are close in the seriation.

| D | 1 | 2 | 3 | 4 | 5 |
|---|-------|------|------|------|------|
| 1 | 0,00 | | | | |
| 2 | 2,61 | 0,00 | | | |
| 3 | 5,33 | 2,73 | 0,00 | | |
| 4 | 7,90 | 5,29 | 2,58 | 0,00 | |
| 5 | 10,55 | 7,94 | 5,22 | 2,65 | 0,00 |



Hierarchical Clusters from Before and After Seriation

- Note seriation (from previous page) cleans up the hierarchy
- In the non-seriated hierarchy we have the green cluster crossing the pink cluster
- Note also that things that belong together tend to be next to each other even without the clustering
- That is, 1 is next to 2, 3 is next to 4, (3,4) is next to 5 etc
- If we plotted this as a heatmap, after seriation we could see where we should make our clusters without running a clustering algorithm
- This will become more obvious when we do our real example



The Setup...

- Last semester in Text Mining, Chris, John, Matt, Mihir, Parvati, Yousuf and others created a corpus of Linked-in docs, web pages, press releases, google scholar articles etc. on a topic of their choice
- They then removed stopwords, lemmatized, etc. their corpus
- After the semester I put all 1,292 documents into one giant corpus to test whether k-means clustering would be smart enough to separate each set into it's proper topic
- Results are on the next page

Results

| | Cluster # | True Positive (Belongs) | False Positive (Doesn't Belong) | False Negative (Should be in) | F Score |
|---------|-----------|----------------------------|------------------------------------|----------------------------------|---------|
| Chris | 9 | 191 | 40 | 0 | 0.905 |
| Eric | 3 | 153 | 12 | 0 | 0.962 |
| John | 0 | 111 | 8 | 0 | 0.965 |
| Matt | 4 | 78 | 11 | 19 | 0.839 |
| Mihir | 7 | 97 | 31 | 2 | 0.855 |
| Parvati | 1 | 43 | 2 | 190 | 0.309 |
| Stephen | 5 | 98 | 34 | 1 | 0.848 |
| Tim | 6 | 100 | 12 | 0 | 0.943 |
| Tony | 2 | 146 | 16 | 0 | 0.948 |
| Yousuf | 8 | 63 | 46 | 0 | 0.733 |
| all | all | 1080 | 212 | 212 | 0.836 |

| | word1 | word2 | word3 | word4 | word5 | word6 | word7 | word8 | word9 | word10 |
|-----------|----------------|------------|-------------|-------------|---------------|---------------|------------|--------------|-------------|----------|
| Cluster 0 | limb | prosthetic | advanced | arm | artificial | prosthesis | hand | amputee | control | darpa |
| Cluster 1 | biotechnology | university | verified | game | email | edu | video | information | engineering | science |
| Cluster 2 | text | mining | information | document | analysis | literature | gene | tool | biomedical | drug |
| Cluster 3 | trump | donald | clinton | woman | debate | campaign | republican | presidential | president | people |
| Cluster 4 | biophysics | chemistry | journal | biophysical | physical | molecular | et | cell | al | patent |
| Cluster 5 | quantum | compute | computer | physic | qubits | university | institute | science | state | waterloo |
| Cluster 6 | airplane | wing | plane | flight | fly | mode | design | pilot | aircraft | jan |
| Cluster 7 | machine | learning | learn | proceeding | international | conference | cloud | google | api | workshop |
| Cluster 8 | nanotechnology | nature | nanoscience | ibm | ieee | nanomaterials | tj | nano | center | watson |
| Cluster 9 | game | video | play | player | violent | study | time | world | child | gaming |

Results (II)

- Note that k-means does the best on cluster 0 (prosthetics) because the topic has nothing to do with the others
- It also does well on the politics, airplane, and gaming clusters
- It inexplicably puts gaming in the biotechnology cluster though
- Overall it does a decent job, and we could make it better by making things like email, et al, jan, edu, etc. into stopwords

Seriation of Documents

- The question is whether we could do a decent job of clustering the documents with seriation?
- How?
- Take entire corpus; do TF/IDF to get a list of 1000 or 2000 most important words
- (A human edited list would work better, but we don't have time)
- Each document is then a 2000 dimension sparse vector with term frequencies
- Seriation puts closest vectors together which ought to work as a clustering method

Step 1: Read in Corpus

```
In [174]: from sklearn.datasets import load_files
...:
...: ## folder containing 1292 lemmatized text files (stopwords removed)
...: ## about 100 files each under names Chris, Parvati, etc.
...: p="E:\Rowan\Classes\TextMining\hw3CaseStudy\LemmatizewordContainer"
...:
...:
...:
...: ## Load all files with one command
...: dataset=load_files(p, description=None, categories=None,
...: load_content=True, shuffle=True, encoding=None,
...: decode_error='strict', random_state=0)
...:
...:
...:
...: print("%d documents" % len(dataset.data))
...: print("%d categories" % len(dataset.target_names))
...: print(dataset.target_names[0:9])

1292 documents
10 categories
['Chris', 'Eric', 'John', 'Matt', 'Mihir', 'Parvati', 'Stephen', 'Tim', 'Tony']
```

- A great thing about sklearn is that if you keep your corpus as a bunch of text files in a hierarchical directory, `load_files` will load all of the files and keep the directory structure; all in one line

Step 2: Make List of all words; sort by freq.

```
In [175]: wordDict = {}
...: for i in range(len(dataset.data)):
...:     words = dataset.data[i].split()
...:     for word in words:
...:         if wordDict.has_key(word):
...:             wordDict[word]+=1
...:         else:
...:             wordDict.setdefault(word,1)
...:
...: import operator
...: wordDist = [sorted(wordDict.iteritems(), key=operator.itemgetter(1),reverse=True)]
...: for s in wordDist:
...:     for (v,k) in s:
...:         if (k > 1000):
...:             print(v+" "+str(k))
game 4623
trump 3727
video 1904
penn 1540
anonymous 1268
time 1190
quantum 1052
airplane 1004
```

- A Python dictionary is useful for keeping track of the occurrences of each word
- Don't know why 'penn' and 'anonymous' appear so frequently.

Step 3: Make List of Top 2000 Words

```
...: ## kill additional stopwords; we could go crazy, but we'll just kill these 2.
In [176]: wordDict['penn']=0
...: wordDict['anonymous']=0
...:
...:
...: ## put top 2000 words in a list
...: ## we won't bother doing IDF
...: count=0
...: list=[]
...: for s in wordDict:
...:     for (v,k) in s:
...:         if (count < 2000):
...:             list.append(v)
...:             count+=1
```

- 2000 is just arbitrary. I don't know if it would work better with 100, 500, 1000, 10000.
- Note also, the first word and 2000th word have the same weight here. Don't know if this will be a problem

Step 4: Write out each document as a sparse vector

```
In [177]: f5=open('vecs2c.txt','w')
...: papDict = {}
...: for i in range(len(dataset.data)):
...:     ## build a dictionary for each paper
...:     words = dataset.data[i].split()
...:     for word in words:
...:         if papDict.has_key(word):
...:             papDict[word]+=1
...:         else:
...:             papDict.setdefault(word,1)
...:     ## now build a vector containing frequencies
...:     vecc=[]
...:     for w in list:
...:         if papDict.has_key(w):
...:             vecc.append(papDict[w])
...:         else:
...:             vecc.append(0)
...:     ## print vector out
...:     f5.write(dataset.target_names[labels[i]]+"_"+str(i))
...:     for t in vecc:
...:         f5.write(", "+str(t))
...:         f5.write("\n")
...:     papDict={}
...: f5.close()
```

- Use a dictionary for each document. Then check to see if any words from the document match the top 2000 words

Snapshot of output
(Total of 1292 vectors; each has 2000 elements)

[illegible]

Go to R

- Go to Seriation2000.pdf

Unrelated Seriation Demo in R

- This one has nothing to do with clustering.
- But it might be the coolest thing you've seen in a while
- Go to [SeriationLouvre.pdf](#)